

# AutoML FraudShield – SRS + TDD + Deployment Manifest (v1)

## 0. Project Overview

**Project Name:** AutoML FraudShield

**Goal:** Build a production-grade, continuously learning fraud detection system that:

- Ingests transactional data (batch + streaming)
- Maintains an online/offline feature store
- Trains and evaluates models via a reproducible ML pipeline
- Registers, versions, and deploys models via a model registry
- Monitors performance & drift and triggers retraining when needed
- Exposes a low-latency scoring API for real-time predictions

**Primary Audience:**

- Hiring managers & teams evaluating you as an AI/ML Engineer, MLOps Engineer, or AI Architect.
- Secondary: Fintech / fraud / risk stakeholders who want to see a full ML lifecycle.

**Target Platform (v1):** Google Cloud Platform

- **Data & Storage:** BigQuery, Cloud Storage
- **Feature Store:** Vertex AI Feature Store
- **Pipeline Orchestration:** Vertex AI Pipelines (orchestrated via Python SDK)
- **Training & Experiments:** Vertex AI Custom Training + Vertex Experiments
- **Model Registry:** Vertex AI Model Registry

- **Serving:** FastAPI on Cloud Run (online scoring)
  - **Monitoring & Drift:** BigQuery + scheduled pipeline + dashboards (Looker / custom)
  - **IaC:** Terraform
- 

## 1. Scope & Requirements (SRS)

### 1.1 In-Scope (v1)

#### 1. Data ingestion (batch)

- Load historical transactions from CSV/Parquet files into BigQuery.
- Apply schema validation and basic quality checks.

#### 2. Feature engineering & feature store

- Define a small but realistic feature set for fraud detection:
  - Transaction-level features (amount, time, merchant, channel).
  - Customer-level aggregates (rolling sum/count, average ticket, recency).
  - Device/geo features (country, IP risk score – stubbed or derived).
- Materialize features into **Vertex AI Feature Store**:
  - Offline store (BigQuery-backed).
  - Online store (low latency feature reads for scoring API).

#### 3. Model training & evaluation

- Implement a reproducible pipeline that:
  - Pulls labeled data from the offline feature store / BigQuery.
  - Splits into train/validation/test with time-aware split.

- Trains at least 2–3 models (e.g., XGBoost, LightGBM, a simple DNN).
- Logs metrics and artifacts (ROC AUC, PR AUC, confusion metrics, calibration).
- Stores metrics in Vertex Experiments.

#### 4. Hyperparameter optimization (HPO)

- Run a small hyperparameter search (grid or Bayesian) for the primary model.
- Automatically select the best model based on ROC AUC (primary) and PR AUC (secondary).

#### 5. Model registry and deployment

- Register the winning model in **Vertex Model Registry** with:
  - Model version
  - Training data reference
  - Metrics
- Deploy the best model to:
  - A **Vertex Endpoint** for internal use (optional but nice), and
  - A **FastAPI scoring service on Cloud Run** that calls:
    - Vertex Endpoint **or**
    - A local model artifact (configurable for portability).

#### 6. Online prediction API

- FastAPI / Cloud Run service:
  - **POST /v1/score** that accepts a JSON payload of transaction attributes.
  - Calls feature store for additional features given entity IDs.
  - Calls model (Vertex Endpoint or local artifact) and returns:

- Fraud probability (0–1)
- Risk band (e.g., Low/Med/High)
- Model version used
- Optional explanation stub (e.g., SHAP in v2).

## 7. Monitoring & drift

- Log all predictions + ground truth labels (when they arrive) to BigQuery.
- Implement a **scheduled Vertex Pipeline** or Cloud Scheduler job that:
  - Computes performance over recent time windows.
  - Computes input distribution drift vs. training baseline.
  - Emits a “retrain recommended” signal when thresholds are breached.

## 8. Automatic retraining (triggered)

- Expose a retraining pipeline entrypoint that:
  - Can be run manually (CLI / SDK call).
  - Can be triggered by a Cloud Scheduler job when:
    - Data volume threshold met (e.g., N new labeled days), or
    - Drift metrics exceed threshold.

## 1.2 Out-of-Scope (v1)

- Complex multi-tenant isolation and billing.
- Formal approval workflows / human-in-the-loop UI.
- Real production-grade label ingestion from external systems (we'll mock labels).
- Fully robust security hardening & compliance (we'll follow best practices but not formal audits).

## **1.3 Functional Requirements (FR)**

**FR-1** – The system must train at least one baseline model and one tuned model and store results.

**FR-2** – The pipeline must log metrics including ROC AUC, PR AUC, and confusion matrix.

**FR-3** – The system must register the best model to a model registry with metadata.

**FR-4** – The scoring API must handle at least 50–100 RPS at p95 latency < 200 ms (with warm instances).

**FR-5** – Predictions must include model version for traceability.

**FR-6** – The drift job must compute at least one feature distribution metric (e.g., PSI, KL divergence, or simple mean/std diff) vs. the training window.

**FR-7** – The retraining pipeline must be runnable via CLI / SDK without manual reconfiguration.

## **1.4 Non-Functional Requirements (NFR)**

### **NFR-1 – Reproducibility:**

All training and evaluation must be traceable and reproducible from code + configs.

### **NFR-2 – Observability:**

- Key stages emit logs + metrics.
- Training runs are visible in Vertex Experiments or MLflow-like UI.

### **NFR-3 – Configurability:**

- Environment configuration (dev/stage/prod) controlled via config files + Terraform variables.

### **NFR-4 – Cost awareness:**

- Pipelines should be designed to run on modest resources (e.g., small custom training cluster).

---

## **2. High-Level Architecture (TDD – System View)**

### **2.1 Components**

1. Data ingestion & prep

- Scripts to load CSV/Parquet into BigQuery.
- Data validation using `great_expectations` or `pandera` (optional).

## 2. Feature engineering service

- Python module to generate offline features.
- Vertex AI Feature Store entity types:
  - `customers`
  - `cards`
  - `merchants`
- Feature groups for aggregates (e.g., 7-day transaction count).

## 3. Training pipeline (Vertex AI Pipelines)

- Steps:
  - Load training data (BigQuery → DataFrame).
  - Join with features from Feature Store (offline).
  - Train baseline and tuned models.
  - Evaluate and compare.
  - Register best model.
  - Optionally deploy to Vertex Endpoint.

## 4. Model registry

- Vertex Model Registry entries with:
  - `model_id`, `version`, `metrics`, `train_data_ref`, `created_at`.

## 5. Scoring API

- FastAPI application:
  - POST /v1/score
- Reads online features from Feature Store given `customer_id` & `card_id`.
- Calls Vertex Endpoint or local model.

## 6. Monitoring & retraining

- BigQuery table: `fraudshield.predictions_log`
- Scheduled job:
  - Pulls last N days predictions + labels.
  - Computes performance + drift metrics.
  - Triggers training pipeline when conditions met.

## 2.2 Data Flow (Conceptual)

### 1. Initial setup

- Load historical labeled transactions → BigQuery.
- Materialize initial features into Feature Store.

### 2. Model training

- Pipeline runs:
  - Query training dataset (transactions + features).
  - Train & evaluate.
  - Register + deploy best model.

### 3. Online scoring

- Client → FastAPI `/score` with transaction payload.

- Service:
  - Fetch features (online FS) for entities.
  - Call model.
  - Log prediction + payload to BigQuery.

#### 4. Monitoring & retrain

- Scheduled job computes performance & drift.
  - If thresholds exceeded → run training pipeline again.
  - New best model registered & deployed.
- 

## 3. Detailed Design (TDD – Component-Level)

### 3.1 Data Model (Simplified)

BigQuery Table `fraudshield.transactions`

- `transaction_id` (STRING, PK)
- `customer_id` (STRING)
- `card_id` (STRING)
- `merchant_id` (STRING)
- `timestamp` (TIMESTAMP)
- `amount` (FLOAT64)
- `currency` (STRING)
- `channel` (STRING) // e.g., POS, eCommerce, ATM

- `is_fraud` (BOOL)

## Feature Store Entity Types

### 1. `customers`

- key: `customer_id`
- features (examples):
  - `customer_txn_count_7d` (INT64)
  - `customer_txn_amount_sum_7d` (FLOAT64)
  - `customer_avg_ticket_30d` (FLOAT64)
  - `customer_geo_entropy_30d` (FLOAT64 – optional stub)

### 2. `cards`

- key: `card_id`
- features:
  - `card_txn_count_7d`
  - `card_txn_amount_sum_7d`

### 3. `merchants`

- key: `merchant_id`
- features:
  - `merchant_txn_count_30d`
  - `merchant_chargeback_rate_90d` (stuffed or simulated)

BigQuery Table `fraudshield.predictions_log`

- `prediction_id` (STRING)
- `transaction_id` (STRING)
- `customer_id` (STRING)
- `card_id` (STRING)
- `timestamp` (TIMESTAMP)
- `score` (FLOAT64)
- `risk_band` (STRING)
- `model_version` (STRING)
- `is_fraud` (BOOL, nullable)
- `logged_at` (TIMESTAMP)

## 3.2 Training Pipeline Steps (Pseudo-Implementation)

### Step 1 – Load training data

- SQL query over `transactions` with filters (date range, not null label).
- Option to downsample majority class (non-fraud) or use class weights.

### Step 2 – Join with features

- Call Feature Store offline API to join entity features.
- Produce a training DataFrame with:
  - Raw features (amount, channel, etc.)
  - Engineered aggregates (from FS)
  - Label (`is_fraud`)

## **Step 3 – Train baseline & tuned models**

- Baseline: XGBoost with default-ish parameters.
- Tuned: HPO with Optuna or Vertex Hyperparameter Tuning.

## **Step 4 – Evaluation**

- Time-aware split: earliest 70% for train, next 15% for valid, last 15% for test.
- Compute metrics:
  - ROC AUC, PR AUC
  - F1 at a chosen threshold
  - Confusion matrix
- Log to Vertex Experiments.

## **Step 5 – Model selection & registry**

- Compare candidate models by ROC AUC (primary) and PR AUC (tie-break).
- Push winning model to Model Registry with:
  - `metrics.json`
  - `training_pipeline_id`
  - `data_ref` (e.g., BQ snapshot + FS version).

## **Step 6 – Deployment**

- Option 1: Deploy to Vertex Endpoint (standard deployment).
- Option 2: Export model artifact (e.g., `model1.pk1`) to GCS where Cloud Run can mount or download.

### **3.3 Scoring API Design**

**Endpoint:** POST /v1/score

**Request (example):**

```
{  
    "transaction_id": "tx_123",  
    "customer_id": "cust_456",  
    "card_id": "card_789",  
    "merchant_id": "m_001",  
    "timestamp": "2025-01-15T12:34:56Z",  
    "amount": 125.40,  
    "currency": "USD",  
    "channel": "ECOM"  
}
```

**Steps:**

1. Validate payload.
2. Read online features from Feature Store:
  - `customers` by `customer_id`
  - `cards` by `card_id`
  - `merchants` by `merchant_id`
3. Assemble feature vector for the model.
4. Call Vertex Endpoint or local model artifact.
5. Map score to risk band, e.g.:
  - `score < 0.2` → LOW
  - `0.2 ≤ score < 0.6` → MEDIUM
  - `score ≥ 0.6` → HIGH
6. Write log row to `predictions_log`.

**Response (example):**

```
{  
  "transaction_id": "tx_123",  
  "score": 0.82,  
  "risk_band": "HIGH",  
  "model_version": "fraudshield-xgb-v5",  
  "explanations": null  
}
```

---

## 4. Deployment Manifest (DM – Infra & Environments)

### 4.1 Environments

- **dev**: single project, low quotas, cheap resources, debug logging.
- **prod**: separate GCP project, stricter IAM, higher quotas, monitoring.

### 4.2 GCP Resources (per project)

- **BigQuery**
  - Dataset: `fraudshield`
  - Tables: `transactions`, `predictions_log`, plus intermediate views.
- **Cloud Storage**
  - `gs://fraudshield-raw-{env}` (raw CSV/Parquet)
  - `gs://fraudshield-artifacts-{env}` (models, configs)
  - `gs://fraudshield-pipeline-logs-{env}`
- **Vertex AI**
  - Feature Store:

- Entity types: `customers`, `cards`, `merchants`
- Pipelines:
  - `fraudshield_training_pipeline`
  - `fraudshield_monitoring_pipeline`
- Model Registry:
  - Artifact name prefix: `fraudshield-`
- (Optional) Online Endpoint:
  - `fraudshield-online-endpoint`
- **Cloud Run**
  - Service: `fraudshield-api-{env}` (FastAPI scoring)
  - Min instances: 1 (for low cold start).
- **Cloud Scheduler**
  - Job: `fraudshield-monitoring-daily` → triggers monitoring pipeline.
  - Job: `fraudshield-retraining-weekly` (optional; may instead be triggered by monitoring).
- **Logging/Monitoring**
  - Cloud Logging & Monitoring dashboards built on:
    - API latency & error rates
    - Training runs
    - Drift metrics (visualized from BigQuery tables)

## 4.3 Terraform Layout

- `infra/terraform/envs/dev/main.tf` – Dev project + resources.
  - `infra/terraform/envs/prod/main.tf` – Prod project + resources.
  - Variables for project ID, region, bucket names, etc.
- 

## 5. Repository / Directory Structure

Proposed repo layout:

```
fraudshield/
├── README.md
├── pyproject.toml          # or setup.cfg/requirements.txt
├── docs/
│   ├── FraudShield_SRS_TDD_DM_v1.md
│   └── architecture_diagrams/  # (mermaid or PNG)
└── infra/
    └── terraform/
        ├── modules/
        │   ├── bigquery/
        │   ├── feature_store/
        │   ├── vertex_pipelines/
        │   ├── cloud_run_api/
        │   └── monitoring/
        └── envs/
            ├── dev/
            │   ├── main.tf
            │   └── variables.tf
            └── prod/
                ├── main.tf
                └── variables.tf
└── data/
    ├── sample_transactions.csv
    └── schemas/
        ├── transactions_schema.json
        └── predictions_log_schema.json
└── features/
```

```
|   └── feature_definitions.py      # high-level feature sets
|   └── build_offline_features.py
|   └── register_feature_store.py
|   └── tests/
|       └── test_feature_defs.py
└── pipelines/
    ├── training/
    |   ├── pipeline_definition.py # Vertex AI Pipelines DAG
    |   └── components/
    |       ├── load_data_component.py
    |       ├── build_dataset_component.py
    |       ├── train_model_component.py
    |       ├── evaluate_model_component.py
    |       └── register_model_component.py
    |   └── launch_training_pipeline.py
    ├── monitoring/
    |   ├── drift_monitoring_pipeline.py
    |   └── launch_monitoring_pipeline.py
    └── utils/
        ├── metrics.py
        └── bq_utils.py
└── api/
    ├── app/
    |   ├── main.py          # FastAPI app
    |   ├── config.py
    |   ├── schemas.py       # Pydantic models
    |   ├── routers/
    |   |   └── scoring.py
    |   ├── services/
    |   |   ├── feature_store_client.py
    |   |   ├── model_client.py # Vertex Endpoint or local model
    |   |   └── logging_client.py
    |   └── tests/
    |       ├── test_scoring_endpoint.py
    |       └── test_feature_fetch.py
    └── Dockerfile
└── models/
```

```
|   └── train_model.py          # direct training script (for local  
dev)  
|   └── model_utils.py  
|   └── artifacts/            # local model artifacts for dev  
└── notebooks/  
    ├── 01_exploration.ipynb  
    ├── 02_feature_prototyping.ipynb  
    └── 03_model_sanity_check.ipynb
```

---

## 6. Roadmap / Future Enhancements

- **v2:** Add SHAP-based explanations to `/score` responses.
- **v2:** Add multi-tenant support via `tenant_id` dimension.
- **v2:** Add UI dashboard (Streamlit) for:
  - Monitoring metrics
  - Triggering re-trains
  - Browsing model versions
- **v3:** Replace Vertex Feature Store with Feast for portability (optional).